

Assignment 8.3

1. Setting HIVE properties in order to work with Transactions in HIVE

```
hive> set hive.support.concurrency = true;
hive> set hive.enforce.bucketing = true;
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> set hive.txn.manager = org.apache.hadoop.hive.ql.lockmgr.DbTxnManager;
hive> set hive.compactor.initiator.on = true;
hive> set hive.compactor.worker.threads=1;
hive> █
```

2. Created table that supports HIVE transactions

```
hive> CREATE TABLE employee(
>   emp_id int,
>   emp_name string,
>   emp_salary int,
>   emp_department string,
>   unit int
> )
> clustered by (emp_id) into 5 buckets stored as orc TBLPROPERTIES('transactional'='true');
OK
Time taken: 0.376 seconds
hive> █
```

3. Insert data into HIVE table:

```
hive> INSERT INTO table employee values (101,'Amitabh',256,'Finance',1),(102,'Shahrukh',78,'IT_Dept',2),(103,'Akshay',110,'HR',3),
> (104,'Anubhav',50,'Network Team',4),(105,'Pawan',250,'Admin',5),(106,'Aamir',25,'Finance',1),(107,'Salman',175,'IT_Dept',2),(108,'Ranbir',142,'HR',3),(109,'
Katrina',100,'Network Team',4),(110,'Priyanka',222,'Admin',5),(111,'Tushar',500,'Finance',1),(112,'Ajay',58,'IT_Dept',2),(113,'Jubeen',100,'Finance',1),(114,'Madh
uri',200,'IT_Dept',2);
Query ID = acadgild_20171120210101_cc63d109-5ecc-4de4-b64c-b73a01a0c0a5
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1511187342194_0001, Tracking URL = http://localhost:8088/proxy/application_1511187342194_0001/
Kill Command = /home/acadgild/hadoop-2.6.0/bin/hadoop job -kill job_1511187342194_0001
Hadoop job Information for Stage-1: number of mappers: 1; number of reducers: 0
2017-11-20 21:02:23,990 Stage-1 map = 0%, reduce = 0%
2017-11-20 21:02:37,420 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.96 sec
MapReduce Total cumulative CPU time: 2 seconds 960 msec
Ended Job = job_1511187342194_0001
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://localhost:9000/tmp/hive/acadgild/59ebbe3f-29b8-4435-b816-7df81cab3cdd/hive_2017-11-20_21-01-56_712_7984791112795110294-1/-ext-10000
Loading data to table sumona.employee
Table sumona.employee stats: [numFiles=1, numRows=14, totalSize=729, rawDataSize=2688]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 2.96 sec HDFS Read: 622 HDFS Write: 802 SUCCESS
Total MapReduce CPU Time Spent: 2 seconds 960 msec
OK
Time taken: 44.174 seconds
hive> █
```

```
hive> SELECT * FROM employee;
OK
101    Amitabh 256      Finance 1
102    Shahrukh 78      IT_Dept 2
103    Akshay 110      HR      3
104    Anubhav 50      Network_Team 4
105    Pawan 250      Admin 5
106    Aamir 25      Finance 1
107    Salman 175     IT_Dept 2
108    Ranbir 142     HR      3
109    Katrina 100     Network_Team 4
110    Priyanka 222     Admin 5
111    Tushar 500     Finance 1
112    Ajay 58      IT_Dept 2
113    Jubeen 100     Finance 1
114    Madhuri 200    IT_Dept 2
Time taken: 0.221 seconds, Fetched: 14 row(s)
hive>
```

4. Updating the values in the table:

UPDATE employee set emp_id = 108 where emp_id = 112;

The above command is used to update a row in HIVE table.

```
hive> UPDATE employee set emp_id = 108 where emp_id = 112;
FAILED: SemanticException [Error 10302]: Updating values of bucketing columns is not supported. Column emp_id.
```

UPDATE college set clg_name = 'IIT' where clg_id = 6;

The above command is used to update on NON bucketed column

```

hive> UPDATE employee set emp name = 'Anushka' where emp id = 105;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20171121122113_90f19e37-b1ef-4f1a-b889-97c30ff4fe07
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1511246161557_0002, Tracking URL = http://localhost:8088/proxy/application_1511246161557_0002/
Kill Command = /home/acadgild/hadoop-2.7.2/bin/hadoop job -kill job_1511246161557_0002
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2017-11-21 12:21:22,561 Stage-1 map = 0%, reduce = 0%
2017-11-21 12:21:47,201 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 1.88 sec
2017-11-21 12:21:49,763 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 3.68 sec
2017-11-21 12:21:52,519 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 7.1 sec
2017-11-21 12:21:53,805 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 8.89 sec
2017-11-21 12:21:55,055 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 9.09 sec
2017-11-21 12:22:12,743 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 10.16 sec
2017-11-21 12:22:14,006 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 11.21 sec
2017-11-21 12:22:15,463 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 12.54 sec
2017-11-21 12:22:16,643 Stage-1 map = 100%, reduce = 73%, Cumulative CPU 14.46 sec
2017-11-21 12:22:18,071 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 16.29 sec
MapReduce Total cumulative CPU time: 16 seconds 290 msec
Ended Job = job_1511246161557_0002
Loading data to table sumona.employee
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 16.29 sec HDFS Read: 62182 HDFS Write: 1113 SUCCESS
Total MapReduce CPU Time Spent: 16 seconds 290 msec
OK
Time taken: 66.574 seconds
hive>

```

Now, if we compare the initial table and the table after updating the row, you will notice that the row 105 has been updated with the name Anushka.

INITIAL TABLE

```

hive> SELECT * FROM employee;
OK
101  Amitabh 256  Finance 1
102  Shahrukh 78  IT_Dept 2
103  Akshay 110  HR 3
104  Anubhav 50  Network_Team 4
105  Pawan 250  Admin 5
106  Aamir 25  Finance 1
107  Salman 175  IT_Dept 2
108  Ranbir 142  HR 3
109  Katrina 100  Network_Team 4
110  Priyanka 222  Admin 5
111  Tushar 500  Finance 1
112  Ajay 58  IT_Dept 2
113  Jubeen 100  Finance 1
114  Madhuri 200  IT_Dept 2
Time taken: 0.221 seconds, Fetched: 14 row(s)
hive>

```

TABLE AFTER UPDATING THE ROW

```

hive> SELECT * FROM employee;
OK
105  Anushka 250  Admin 5
110  Priyanka 222  Admin 5
106  Aamir 25  Finance 1
101  Amitabh 256  Finance 1
111  Tushar 500  Finance 1
107  Salman 175  IT_Dept 2
112  Ajay 58  IT_Dept 2
102  Shahrukh 78  IT_Dept 2
113  Jubeen 100  Finance 1
103  Akshay 110  HR 3
108  Ranbir 142  HR 3
109  Katrina 100  Network_Team 4
104  Anubhav 50  Network_Team 4
114  Madhuri 200  IT_Dept 2
Time taken: 0.184 seconds, Fetched: 14 row(s)
hive>

```

5. Deleting a Row from HIVE table:

```
hive> delete from employee where emp_id= 107;
WARNING: Hive-on-MR is deprecated in Hive 2 and may not be available in the future versions. Consider using a different execution engine (i.e. spark, tez) or using Hive 1.X releases.
Query ID = acadgild_20171121123313_50eaf3bc-cb02-4c9a-929a-c83e6bed4881
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 5
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1511246161557_0003, Tracking URL = http://localhost:8088/proxy/application_1511246161557_0003/
Kill Command = /home/acadgild/hadoop-2.7.2/bin/hadoop job -kill job_1511246161557_0003
Hadoop job information for Stage-1: number of mappers: 5; number of reducers: 5
2017-11-21 12:33:21,782 Stage-1 map = 0%, reduce = 0%
2017-11-21 12:33:41,829 Stage-1 map = 20%, reduce = 0%, Cumulative CPU 1.76 sec
2017-11-21 12:33:45,857 Stage-1 map = 40%, reduce = 0%, Cumulative CPU 3.57 sec
2017-11-21 12:33:50,104 Stage-1 map = 60%, reduce = 0%, Cumulative CPU 6.95 sec
2017-11-21 12:33:51,439 Stage-1 map = 80%, reduce = 0%, Cumulative CPU 8.72 sec
2017-11-21 12:33:54,018 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 8.99 sec
2017-11-21 12:34:06,472 Stage-1 map = 100%, reduce = 13%, Cumulative CPU 10.08 sec
2017-11-21 12:34:07,901 Stage-1 map = 100%, reduce = 27%, Cumulative CPU 11.15 sec
2017-11-21 12:34:09,165 Stage-1 map = 100%, reduce = 40%, Cumulative CPU 11.78 sec
2017-11-21 12:34:10,320 Stage-1 map = 100%, reduce = 53%, Cumulative CPU 12.9 sec
2017-11-21 12:34:11,447 Stage-1 map = 100%, reduce = 60%, Cumulative CPU 13.33 sec
2017-11-21 12:34:12,631 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 15.94 sec
MapReduce Total cumulative CPU time: 15 seconds 940 msec
Ended Job = job_1511246161557_0003
Loading data to table sumona.employee
MapReduce Jobs Launched:
Stage-Stage-1: Map: 5 Reduce: 5 Cumulative CPU: 15.94 sec HDFS Read: 57207 HDFS Write: 757 SUCCESS
Total MapReduce CPU Time Spent: 15 seconds 940 msec
OK
Time taken: 61.455 seconds
hive> █
```

Now, if we compare the initial table and the table after deleting the row, you will notice that the row 107 has been deleted.

INITIAL TABLE

```
hive> SELECT * FROM employee;
OK
101  Amitabh 256    Finance 1
102  Shahrukh 78      IT_Dept 2
103  Akshay 110     HR      3
104  Anubhav 50      Network_Team 4
105  Pawan 250     Admin 5
106  Aamir 25      Finance 1
107  Salman 175     IT_Dept 2
108  Ranbir 142     HR      3
109  Katrina 100     Network_Team 4
110  Priyanka 222     Admin 5
111  Tushar 500     Finance 1
112  Ajay 58      IT_Dept 2
113  Jubeen 100     Finance 1
114  Madhuri 200    IT_Dept 2
Time taken: 0.221 seconds, Fetched: 14 row(s)
hive> █
```

TABLE AFTER DLEELTING A ROW

```
hive> SELECT * FROM employee;
OK
105  Anushka 250     Admin 5
110  Priyanka 222     Admin 5
106  Aamir 25      Finance 1
101  Amitabh 256     Finance 1
111  Tushar 500     Finance 1
112  Ajay 58      IT_Dept 2
102  Shahrukh 78      IT_Dept 2
113  Jubeen 100     Finance 1
103  Akshay 110     HR      3
108  Ranbir 142     HR      3
109  Katrina 100     Network_Team 4
104  Anubhav 50      Network_Team 4
114  Madhuri 200    IT_Dept 2
Time taken: 0.142 seconds, Fetched: 13 row(s)
hive> █
```